

---

```

#include "config.h"
#include "mesh_constants_cuda.h"
#include <cuda.h>
#include <cuda_runtime.h>
#include <stdio.h>
__global__ void compute_forces(int nb_blocks_to_compute,
                               int NGLLOB,
                               int* d_ibool,
                               int* d_phase_ispec_inner_elastic, int
num_phase_ispec_elastic,
                               int d_iphase,
                               realw* d_displ, realw* d_veloc, realw* d_accel,
                               realw* d_xix, realw* d_xiy, realw* d_xiz,
                               realw* d_etax, realw* d_etay, realw* d_etaz,
                               realw* d_gammax, realw* d_gammay, realw* d_gammaz,
                               realw* d_hprime_xx,
                               realw* d_hprimewgll_xx,
                               realw* d_wgllwgll_xy, realw* d_wgllwgll_xz, realw*
d_wgllwgll_yz,
                               realw* d_kappav, realw* d_muv,
                               int NSPEC,
                               realw* d_rhostore,
                               realw* wgll_cube,
                               int* maskx,
                               int* maskax,
                               int myrank){

// elastic compute kernel without attenuation
// holds for: ATTENUATION = .false.
//          COMPUTE_AND_STORE_STRAIN = .true. or .false. (true for kernel simulations)

int bx = blockIdx.y*gridDim.x+blockIdx.x;
int tx = threadIdx.x;

const int NGLL3_ALIGN = NGLL3_PADDED;

int K = (tx/NGLL2);
int J = ((tx-K*NGLL2)/NGLLX);
int I = (tx-K*NGLL2-J*NGLLX);

int active, offset;
int iglob = 0;
int working_element;

realw tempx1l, tempx2l, tempx3l, tempy1l, tempy2l, tempy3l, tempz1l, tempz2l, tempz3l;
realw xixl, xiyl, xizl, etaxl, etayl, etazl, gammaxl, gammayl, gammazl, jacobianl;
realw duxdxl, duxdyl, duxdzl, duydxl, duydy, duydzl, duzdxl, duzdyl, duzdzl;
realw duxdxl_plus_duydyl, duxdxl_plus_duzdzl, duydy_plus_duzdzl;
realw duxdyl_plus_duydxl, duzdxl_plus_duxdzl, duzdyl_plus_duydzl;

realw fac1, fac2, fac3, lambdal, mul, lambdalplus2mul, kappal;
realw sigma_xx, sigma_yy, sigma_zz, sigma_xy, sigma_xz, sigma_yz;

realw sum_terms1, sum_terms2, sum_terms3;

// gravity variables
realw sigma_yx, sigma_zx, sigma_zy;

__shared__ realw s_dummyx_loc[NGLL3];

```

```

__shared__ realw s_dummyx_loc[NGLL3];
__shared__ realw s_dummys_loc[NGLL3];

__shared__ realw s_tempx1[NGLL3];
__shared__ realw s_tempx2[NGLL3];
__shared__ realw s_tempx3[NGLL3];

__shared__ realw s_tempx1[NGLL3];
__shared__ realw s_tempx2[NGLL3];
__shared__ realw s_tempx3[NGLL3];

__shared__ realw s_tempz1[NGLL3];
__shared__ realw s_tempz2[NGLL3];
__shared__ realw s_tempz3[NGLL3];

//__shared__ realw sh_hprime_xx[NGLL2];

// use only NGLL^3 = 125 active threads, plus 3 inactive/ghost threads,
// because we used memory padding from NGLL^3 = 125 to 128 to get coalescent memory
// accesses
active = (tx < NGLL3 && bx < nb_blocks_to_compute) ? 1:0;

// printf("\nwe are here: tx = %d\n",tx);
// copy from global memory to shared memory
// each thread writes one of the NGLL^3 = 125 data points
if (active) {

    working_element = d_phase_ispec_inner_elastic[bx + num_phase_ispec_elastic*
(d_iphase-1)]-1;
    iglob = d_ibool[working_element*NGLL3 + tx]-1;
    // debug
    //if( iglob < 0 || iglob >= NGLOB ){ printf("wrong iglob %d\n",iglob); }

    // changing iglob indexing to match fortran row changes fast style
    s_dummyx_loc[tx] = d_displ[iglob*3];
    s_dummys_loc[tx] = d_displ[iglob*3 + 1];
    s_dummys_loc[tx] = d_displ[iglob*3 + 2];
    s_dummyx_loc[tx] = s_dummyx_loc[tx] * maskx[iglob*3] * (-1.0f);
    s_dummys_loc[tx] = s_dummys_loc[tx] * maskx[iglob*3 + 1] * (-1.0f);
    s_dummys_loc[tx] = s_dummys_loc[tx] * maskx[iglob*3 + 2] * (-1.0f);

    // if(!maskx[iglob*3]) printf("maskvaluefalse : %d\n",maskx[iglob*3]);
    // JC JC here we will need to add GPU support for the new C-PML routines
    // if(maskx[iglob*3]) printf("maskvaluetrue : %d\n",maskx[iglob*3]);
    // JC JC here we will need to add GPU support for the new C-PML routines

}

__syncthreads();

if (active) {

    tempx11 = s_dummyx_loc[K*NGLL2+J*NGLLX]*d_hprime_xx[I]
    + s_dummyx_loc[K*NGLL2+J*NGLLX+1]*d_hprime_xx[NGLLX+I]
    + s_dummyx_loc[K*NGLL2+J*NGLLX+2]*d_hprime_xx[2*NGLLX+I]

```

```

+ s_dummyx_loc[K*NGLL2+J*NGLLX+3]*d_hprime_xx[3*NGLLX+I]
+ s_dummyx_loc[K*NGLL2+J*NGLLX+4]*d_hprime_xx[4*NGLLX+I];

tempy1l = s_dummyy_loc[K*NGLL2+J*NGLLX]*d_hprime_xx[I]
+ s_dummyy_loc[K*NGLL2+J*NGLLX+1]*d_hprime_xx[NGLLX+I]
+ s_dummyy_loc[K*NGLL2+J*NGLLX+2]*d_hprime_xx[2*NGLLX+I]
+ s_dummyy_loc[K*NGLL2+J*NGLLX+3]*d_hprime_xx[3*NGLLX+I]
+ s_dummyy_loc[K*NGLL2+J*NGLLX+4]*d_hprime_xx[4*NGLLX+I];

tempz1l = s_dummyz_loc[K*NGLL2+J*NGLLX]*d_hprime_xx[I]
+ s_dummyz_loc[K*NGLL2+J*NGLLX+1]*d_hprime_xx[NGLLX+I]
+ s_dummyz_loc[K*NGLL2+J*NGLLX+2]*d_hprime_xx[2*NGLLX+I]
+ s_dummyz_loc[K*NGLL2+J*NGLLX+3]*d_hprime_xx[3*NGLLX+I]
+ s_dummyz_loc[K*NGLL2+J*NGLLX+4]*d_hprime_xx[4*NGLLX+I];

tempx2l = s_dummyx_loc[K*NGLL2+I]*d_hprime_xx[J]
+ s_dummyx_loc[K*NGLL2+NGLLX+I]*d_hprime_xx[NGLLX+J]
+ s_dummyx_loc[K*NGLL2+2*NGLLX+I]*d_hprime_xx[2*NGLLX+J]
+ s_dummyx_loc[K*NGLL2+3*NGLLX+I]*d_hprime_xx[3*NGLLX+J]
+ s_dummyx_loc[K*NGLL2+4*NGLLX+I]*d_hprime_xx[4*NGLLX+J];

tempy2l = s_dummyy_loc[K*NGLL2+I]*d_hprime_xx[J]
+ s_dummyy_loc[K*NGLL2+NGLLX+I]*d_hprime_xx[NGLLX+J]
+ s_dummyy_loc[K*NGLL2+2*NGLLX+I]*d_hprime_xx[2*NGLLX+J]
+ s_dummyy_loc[K*NGLL2+3*NGLLX+I]*d_hprime_xx[3*NGLLX+J]
+ s_dummyy_loc[K*NGLL2+4*NGLLX+I]*d_hprime_xx[4*NGLLX+J];

tempz2l = s_dummyz_loc[K*NGLL2+I]*d_hprime_xx[J]
+ s_dummyz_loc[K*NGLL2+NGLLX+I]*d_hprime_xx[NGLLX+J]
+ s_dummyz_loc[K*NGLL2+2*NGLLX+I]*d_hprime_xx[2*NGLLX+J]
+ s_dummyz_loc[K*NGLL2+3*NGLLX+I]*d_hprime_xx[3*NGLLX+J]
+ s_dummyz_loc[K*NGLL2+4*NGLLX+I]*d_hprime_xx[4*NGLLX+J];

tempx3l = s_dummyx_loc[J*NGLLX+I]*d_hprime_xx[K]
+ s_dummyx_loc[NGLL2+J*NGLLX+I]*d_hprime_xx[NGLLX+K]
+ s_dummyx_loc[2*NGLL2+J*NGLLX+I]*d_hprime_xx[2*NGLLX+K]
+ s_dummyx_loc[3*NGLL2+J*NGLLX+I]*d_hprime_xx[3*NGLLX+K]
+ s_dummyx_loc[4*NGLL2+J*NGLLX+I]*d_hprime_xx[4*NGLLX+K];

tempy3l = s_dummyy_loc[J*NGLLX+I]*d_hprime_xx[K]
+ s_dummyy_loc[NGLL2+J*NGLLX+I]*d_hprime_xx[NGLLX+K]
+ s_dummyy_loc[2*NGLL2+J*NGLLX+I]*d_hprime_xx[2*NGLLX+K]
+ s_dummyy_loc[3*NGLL2+J*NGLLX+I]*d_hprime_xx[3*NGLLX+K]
+ s_dummyy_loc[4*NGLL2+J*NGLLX+I]*d_hprime_xx[4*NGLLX+K];

tempz3l = s_dummyz_loc[J*NGLLX+I]*d_hprime_xx[K]
+ s_dummyz_loc[NGLL2+J*NGLLX+I]*d_hprime_xx[NGLLX+K]
+ s_dummyz_loc[2*NGLL2+J*NGLLX+I]*d_hprime_xx[2*NGLLX+K]
+ s_dummyz_loc[3*NGLL2+J*NGLLX+I]*d_hprime_xx[3*NGLLX+K]
+ s_dummyz_loc[4*NGLL2+J*NGLLX+I]*d_hprime_xx[4*NGLLX+K];
/* if(myrank == 31 && bx == 0) printf("\ntx = %d, s_dummyx_loc=%f\n",tx,s_dummyx_loc
[tx]);
if(myrank == 31 && bx == 0) printf("\ntx = %d, d_hprime_xx=%f\n",tx,d_hprime_xx[tx]);
if(myrank == 31 && bx == 0) printf("\ntx = %d, tempx1l=%e,tempx2l=%e,tempx3l=%
e,tempy1l=%e\n",tx,tempx1l,tempx2l,tempx3l,tempy1l);
if(myrank == 31 && bx == 0) printf("\nNGLLX = %d, NGLL2=%d\n",NGLLX,NGLL2);
if(myrank == 31 && bx == 0) printf("\ntx = %d, x0=%e,x0c=%e,x1=%e,x1c=%e,x2=%e,x2c=%
e,x3=%e,x3c=%e,x4=%e,x4c=%e\n",tx,
s_dummyx_loc[K*NGLL2+J*NGLLX],d_hprime_xx[I],
s_dummyx_loc[K*NGLL2+J*NGLLX+1],d_hprime_xx[NGLLX+I],
s_dummyx_loc[K*NGLL2+J*NGLLX+2],d_hprime_xx[2*NGLLX+I],

```

```

s_dummyx_loc[K*NGLL2+J*NGLLX+3],d_hprime_xx[3*NGLLX+I],
s_dummyx_loc[K*NGLL2+J*NGLLX+4],d_hprime_xx[4*NGLLX+I]);
*/

// JC JC here we will need to add GPU support for the new C-PML routines

// compute derivatives of ux, uy and uz with respect to x, y and z
offset = working_element*NGLL3_ALIGN + tx;

xixl = d_xix[offset];
xiyl = d_xiy[offset];
xizl = d_xiz[offset];
etaxl = d_etax[offset];
etayl = d_etay[offset];
etazl = d_etaz[offset];
gammaxl = d_gammax[offset];
gammayl = d_gammay[offset];
gammazl = d_gammaz[offset];

duxdxl = xixl*temp1l + etaxl*temp2l + gammaxl*temp3l;
duxdyl = xiyl*temp1l + etayl*temp2l + gammayl*temp3l;
duxdzl = xizl*temp1l + etazl*temp2l + gammazl*temp3l;

duydxl = xixl*temp1l + etaxl*temp2l + gammaxl*temp3l;
duydyl = xiyl*temp1l + etayl*temp2l + gammayl*temp3l;
duydzl = xizl*temp1l + etazl*temp2l + gammazl*temp3l;

duzdxl = xixl*temp1l + etaxl*temp2l + gammaxl*temp3l;
duzdyl = xiyl*temp1l + etayl*temp2l + gammayl*temp3l;
duzdzl = xizl*temp1l + etazl*temp2l + gammazl*temp3l;

//for dbg
/* if(myrank == 31 && iglob == 1) printf("gpu info: %d: duxdxl= %
f",iglob,duxdxl);*/
// JC JC here we will need to add GPU support for the new C-PML routines

// precompute some sums to save CPU time
duxdxl_plus_duydyl = duxdxl + duydyl;
duxdxl_plus_duzdzl = duxdxl + duzdzl;
duydyl_plus_duzdzl = duydyl + duzdzl;
duxdyl_plus_duydxl = duxdyl + duydxl;
duzdxl_plus_duxdzl = duzdxl + duxdzl;
duzdyl_plus_duydzl = duzdyl + duydzl;

// JC JC here we will need to add GPU support for the new C-PML routines

// computes deviatoric strain for kernel calculations

// compute elements with an elastic isotropic rheology
kappal = d_kappav[offset];
mul = d_muv[offset];

// full anisotropic case, stress calculations
// isotropic case

lambdalplus2mul = kappal + 1.3333333333333333 * mul; // 4./3. = 1.3333333
lambdal = lambdalplus2mul - 2.0 * mul;
/*

```

```

        if(tx == 4 && myrank == 31 && bx == 0) printf("\n4__stempx1: %e,%e,%e,%e,%e\n",
lambda, mul, duxdxl, duydy, duzdzl);
        if(tx == 3 && myrank == 31 && bx == 0) printf("\n3__stempx1: %e,%e,%e,%e,%e\n",
lambda, mul, duxdxl, duydy, duzdzl);
    */

    // compute the six components of the stress tensor sigma
    sigma_xx = lambda + 2*mul*duxdxl + lambda*duydy + duzdzl;
    sigma_yy = lambda + 2*mul*duydy + lambda*duxdxl + duzdzl;
    sigma_zz = lambda + 2*mul*duzdzl + lambda*duxdxl + duydy;

    sigma_xy = mul*duxdy + duydx;
    sigma_xz = mul*duzdx + duxdz;
    sigma_yz = mul*duzdy + duydz;

    jacobianl = 1.0f / (xixl*(etayl*gamma_zl - etazl*gamma_y) - xiyl*(etaxl*gamma_zl -
etazl*gamma_x) + xizl*(etaxl*gamma_y - etayl*gamma_x));

/*      if(myrank == 31 && iglob == 1) printf("gpu:jacobian:%f,I,%d,J,%d,K,%d,element,%
d,iphase,%d",jacobianl,I,J,K,working_element,d_iphase);*/
    // define symmetric components (needed for non-symmetric dot product and sigma for
gravity)
    sigma_yx = sigma_xy;
    sigma_zx = sigma_xz;
    sigma_zy = sigma_yz;

    // form dot product with test vector, non-symmetric form
    s_tempx1[tx] = jacobianl * (sigma_xx*xixl + sigma_yx*xiyl + sigma_zx*xizl);
    s_tempy1[tx] = jacobianl * (sigma_xy*xixl + sigma_yy*xiyl + sigma_zy*xizl);
    s_tempz1[tx] = jacobianl * (sigma_xz*xixl + sigma_yz*xiyl + sigma_zz*xizl);

    s_tempx2[tx] = jacobianl * (sigma_xx*etaxl + sigma_yx*etayl + sigma_zx*etazl);
    s_tempy2[tx] = jacobianl * (sigma_xy*etaxl + sigma_yy*etayl + sigma_zy*etazl);
    s_tempz2[tx] = jacobianl * (sigma_xz*etaxl + sigma_yz*etayl + sigma_zz*etazl);

    s_tempx3[tx] = jacobianl * (sigma_xx*gamma_xl + sigma_yx*gamma_y + sigma_zx*gamma_z);
    s_tempy3[tx] = jacobianl * (sigma_xy*gamma_xl + sigma_yy*gamma_y + sigma_zy*gamma_z);
    s_tempz3[tx] = jacobianl * (sigma_xz*gamma_xl + sigma_yz*gamma_y + sigma_zz*gamma_z);

}

// synchronize all the threads (one thread for each of the NGLL grid points of the
// current spectral element) because we need the whole element to be ready in order
// to be able to compute the matrix products along cut planes of the 3D element below
__syncthreads();

// JC JC here we will need to add GPU support for the new C-PML routines

if (active) {

    tempx1l = s_tempx1[K*NGLL2+J*NGLLX]*d_hprimewll_xx[I*NGLLX]
+ s_tempx1[K*NGLL2+J*NGLLX+1]*d_hprimewll_xx[I*NGLLX+1]
+ s_tempx1[K*NGLL2+J*NGLLX+2]*d_hprimewll_xx[I*NGLLX+2]
+ s_tempx1[K*NGLL2+J*NGLLX+3]*d_hprimewll_xx[I*NGLLX+3]
+ s_tempx1[K*NGLL2+J*NGLLX+4]*d_hprimewll_xx[I*NGLLX+4];

    tempy1l = s_tempy1[K*NGLL2+J*NGLLX]*d_hprimewll_xx[I*NGLLX]
+ s_tempy1[K*NGLL2+J*NGLLX+1]*d_hprimewll_xx[I*NGLLX+1]
+ s_tempy1[K*NGLL2+J*NGLLX+2]*d_hprimewll_xx[I*NGLLX+2]
+ s_tempy1[K*NGLL2+J*NGLLX+3]*d_hprimewll_xx[I*NGLLX+3]
+ s_tempy1[K*NGLL2+J*NGLLX+4]*d_hprimewll_xx[I*NGLLX+4];

```

```

tempz1l = s_tempz1[K*NGLL2+J*NGLLX]*d_hprimewgll_xx[I*NGLLX]
+ s_tempz1[K*NGLL2+J*NGLLX+1]*d_hprimewgll_xx[I*NGLLX+1]
+ s_tempz1[K*NGLL2+J*NGLLX+2]*d_hprimewgll_xx[I*NGLLX+2]
+ s_tempz1[K*NGLL2+J*NGLLX+3]*d_hprimewgll_xx[I*NGLLX+3]
+ s_tempz1[K*NGLL2+J*NGLLX+4]*d_hprimewgll_xx[I*NGLLX+4];

tempx2l = s_tempx2[K*NGLL2+I]*d_hprimewgll_xx[J*NGLLX]
+ s_tempx2[K*NGLL2+NGLLX+I]*d_hprimewgll_xx[J*NGLLX+1]
+ s_tempx2[K*NGLL2+2*NGLLX+I]*d_hprimewgll_xx[J*NGLLX+2]
+ s_tempx2[K*NGLL2+3*NGLLX+I]*d_hprimewgll_xx[J*NGLLX+3]
+ s_tempx2[K*NGLL2+4*NGLLX+I]*d_hprimewgll_xx[J*NGLLX+4];

tempy2l = s_tempy2[K*NGLL2+I]*d_hprimewgll_xx[J*NGLLX]
+ s_tempy2[K*NGLL2+NGLLX+I]*d_hprimewgll_xx[J*NGLLX+1]
+ s_tempy2[K*NGLL2+2*NGLLX+I]*d_hprimewgll_xx[J*NGLLX+2]
+ s_tempy2[K*NGLL2+3*NGLLX+I]*d_hprimewgll_xx[J*NGLLX+3]
+ s_tempy2[K*NGLL2+4*NGLLX+I]*d_hprimewgll_xx[J*NGLLX+4];

tempz2l = s_tempz2[K*NGLL2+I]*d_hprimewgll_xx[J*NGLLX]
+ s_tempz2[K*NGLL2+NGLLX+I]*d_hprimewgll_xx[J*NGLLX+1]
+ s_tempz2[K*NGLL2+2*NGLLX+I]*d_hprimewgll_xx[J*NGLLX+2]
+ s_tempz2[K*NGLL2+3*NGLLX+I]*d_hprimewgll_xx[J*NGLLX+3]
+ s_tempz2[K*NGLL2+4*NGLLX+I]*d_hprimewgll_xx[J*NGLLX+4];

tempx3l = s_tempx3[J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX]
+ s_tempx3[NGLL2+J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX+1]
+ s_tempx3[2*NGLL2+J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX+2]
+ s_tempx3[3*NGLL2+J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX+3]
+ s_tempx3[4*NGLL2+J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX+4];

tempy3l = s_tempy3[J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX]
+ s_tempy3[NGLL2+J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX+1]
+ s_tempy3[2*NGLL2+J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX+2]
+ s_tempy3[3*NGLL2+J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX+3]
+ s_tempy3[4*NGLL2+J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX+4];

tempz3l = s_tempz3[J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX]
+ s_tempz3[NGLL2+J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX+1]
+ s_tempz3[2*NGLL2+J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX+2]
+ s_tempz3[3*NGLL2+J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX+3]
+ s_tempz3[4*NGLL2+J*NGLLX+I]*d_hprimewgll_xx[K*NGLLX+4];

fac1 = d_wgllwgll_yz[K*NGLLX+J];
fac2 = d_wgllwgll_xz[K*NGLLX+I];
fac3 = d_wgllwgll_xy[J*NGLLX+I];

sum_terms1 = (fac1*tempx1l + fac2*tempx2l + fac3*tempx3l) * maskax[iglob*3] *
(-1.0f);
sum_terms2 = (fac1*tempy1l + fac2*tempy2l + fac3*tempy3l) * maskax[iglob*3 + 1] *
(-1.0f);
sum_terms3 = (fac1*tempz1l + fac2*tempz2l + fac3*tempz3l) * maskax[iglob*3 + 2] *
(-1.0f);

// adds gravity term
atomicAdd(&d_accel[iglob*3], sum_terms1);
atomicAdd(&d_accel[iglob*3+1], sum_terms2);
atomicAdd(&d_accel[iglob*3+2], sum_terms3);
// if(iglob == 362503 && myrank == 4) printf("\nGPU side: %e daccel: d_accel= %
e, block: %d iphase: %d\n",sum_terms1,d_accel[iglob*3],working_element,d_iphase);

```

---

```

//      if(iglob == 26164  && myrank==11) printf("\nGPU side: %e daccel: d_accel= %
e, block: %d iphase: %d\n",sum_terms1,d_accel[iglob*3],working_element,d_iphase);
//      if(myrank == 11 && working_element == 344) printf("\nGPUint:%d,%d
\n",iglob,maskx[3*iglob]);

/*      if(iglob == 1 && myrank == 31) printf("\ngpu : delta: %f\n",sum_terms1);
      if(iglob == 1 && myrank == 31) printf("\ngpu : fac1 %f tempx1l%f\n, s_tempx1,%
f,tx= %d\n",fac1,tempx1l,s_tempx1[tx],tx);

      if(iglob == 1 && myrank == 31)
      {
        for(int ii = 0;ii<125;ii++) printf("\nstemp%d: = %f\n",ii,s_tempx1[ii]);
      }
      */
} // if(active)
} // kernel_2_noatt_impl()

extern "C"
void FC_FUNC_(compute_forces_fault,
               COMPUTE_FORCES_FAULT)(long* Mesh_pointer,
                                     int* iphase,
                                     realw* deltat,
                                     realw*
CG_d_displ,
                                     realw*
CG_d_accel,
                                     int* maskx,
                                     int* maskax,
                                     int* nspec_outer_elastic,
                                     int* nspec_inner_elastic,
                                     int* myrank )
{
  TRACE("\tcompute_forces_fault");
  // EPIK_TRACER("compute_forces_viscoelastic_cuda");
  //printf("Running compute_forces\n");
  //double start_time = get_time();

  Mesh* mp = (Mesh*)(*Mesh_pointer); // get Mesh from fortran integer wrapper

  int num_elements;

  if( *iphase == 1 )
    num_elements = *nspec_outer_elastic;
  else
    num_elements = *nspec_inner_elastic;

  // checks if anything to do
  if( num_elements == 0 ) return;

  int blocksize = NGLL3_PADDED;

  int num_blocks_x, num_blocks_y;
  get_blocks_xy(num_elements,&num_blocks_x,&num_blocks_y);

  dim3 grid(num_blocks_x,num_blocks_y);
  dim3 threads(blocksize,1,1);

```

```
compute_forces<<<grid,threads>>>(  
    num_elements,  
    mp->NGLOB_AB,  
    mp->d_ibool,  
    mp->d_phase_ispec_inner_elastic,  
    mp->num_phase_ispec_elastic,  
    *iphase,  
    CG_d_displ,    mp->d_veloc,    CG_d_accel,  
    mp->d_xix,      mp->d_xiy,      mp->d_xiz,  
    mp->d_etax,      mp->d_etay,      mp->d_etaz,  
    mp->d_gammax,    mp->d_gammay,    mp->d_gammaz,  
    mp->d_hprime_xx,mp->d_hprimewgll_xx,  
    mp->d_wgllwgll_xy,  
    mp->d_wgllwgll_xz,  
    mp->d_wgllwgll_yz,  
    mp->d_kappav,    mp->d_muv,  
    mp->NSPEC_AB,  
    mp->d_rhostore,  
    mp->d_wgll_cube,  
    maskx,  
    maskax,  
    *myrank);
```

```
}
```